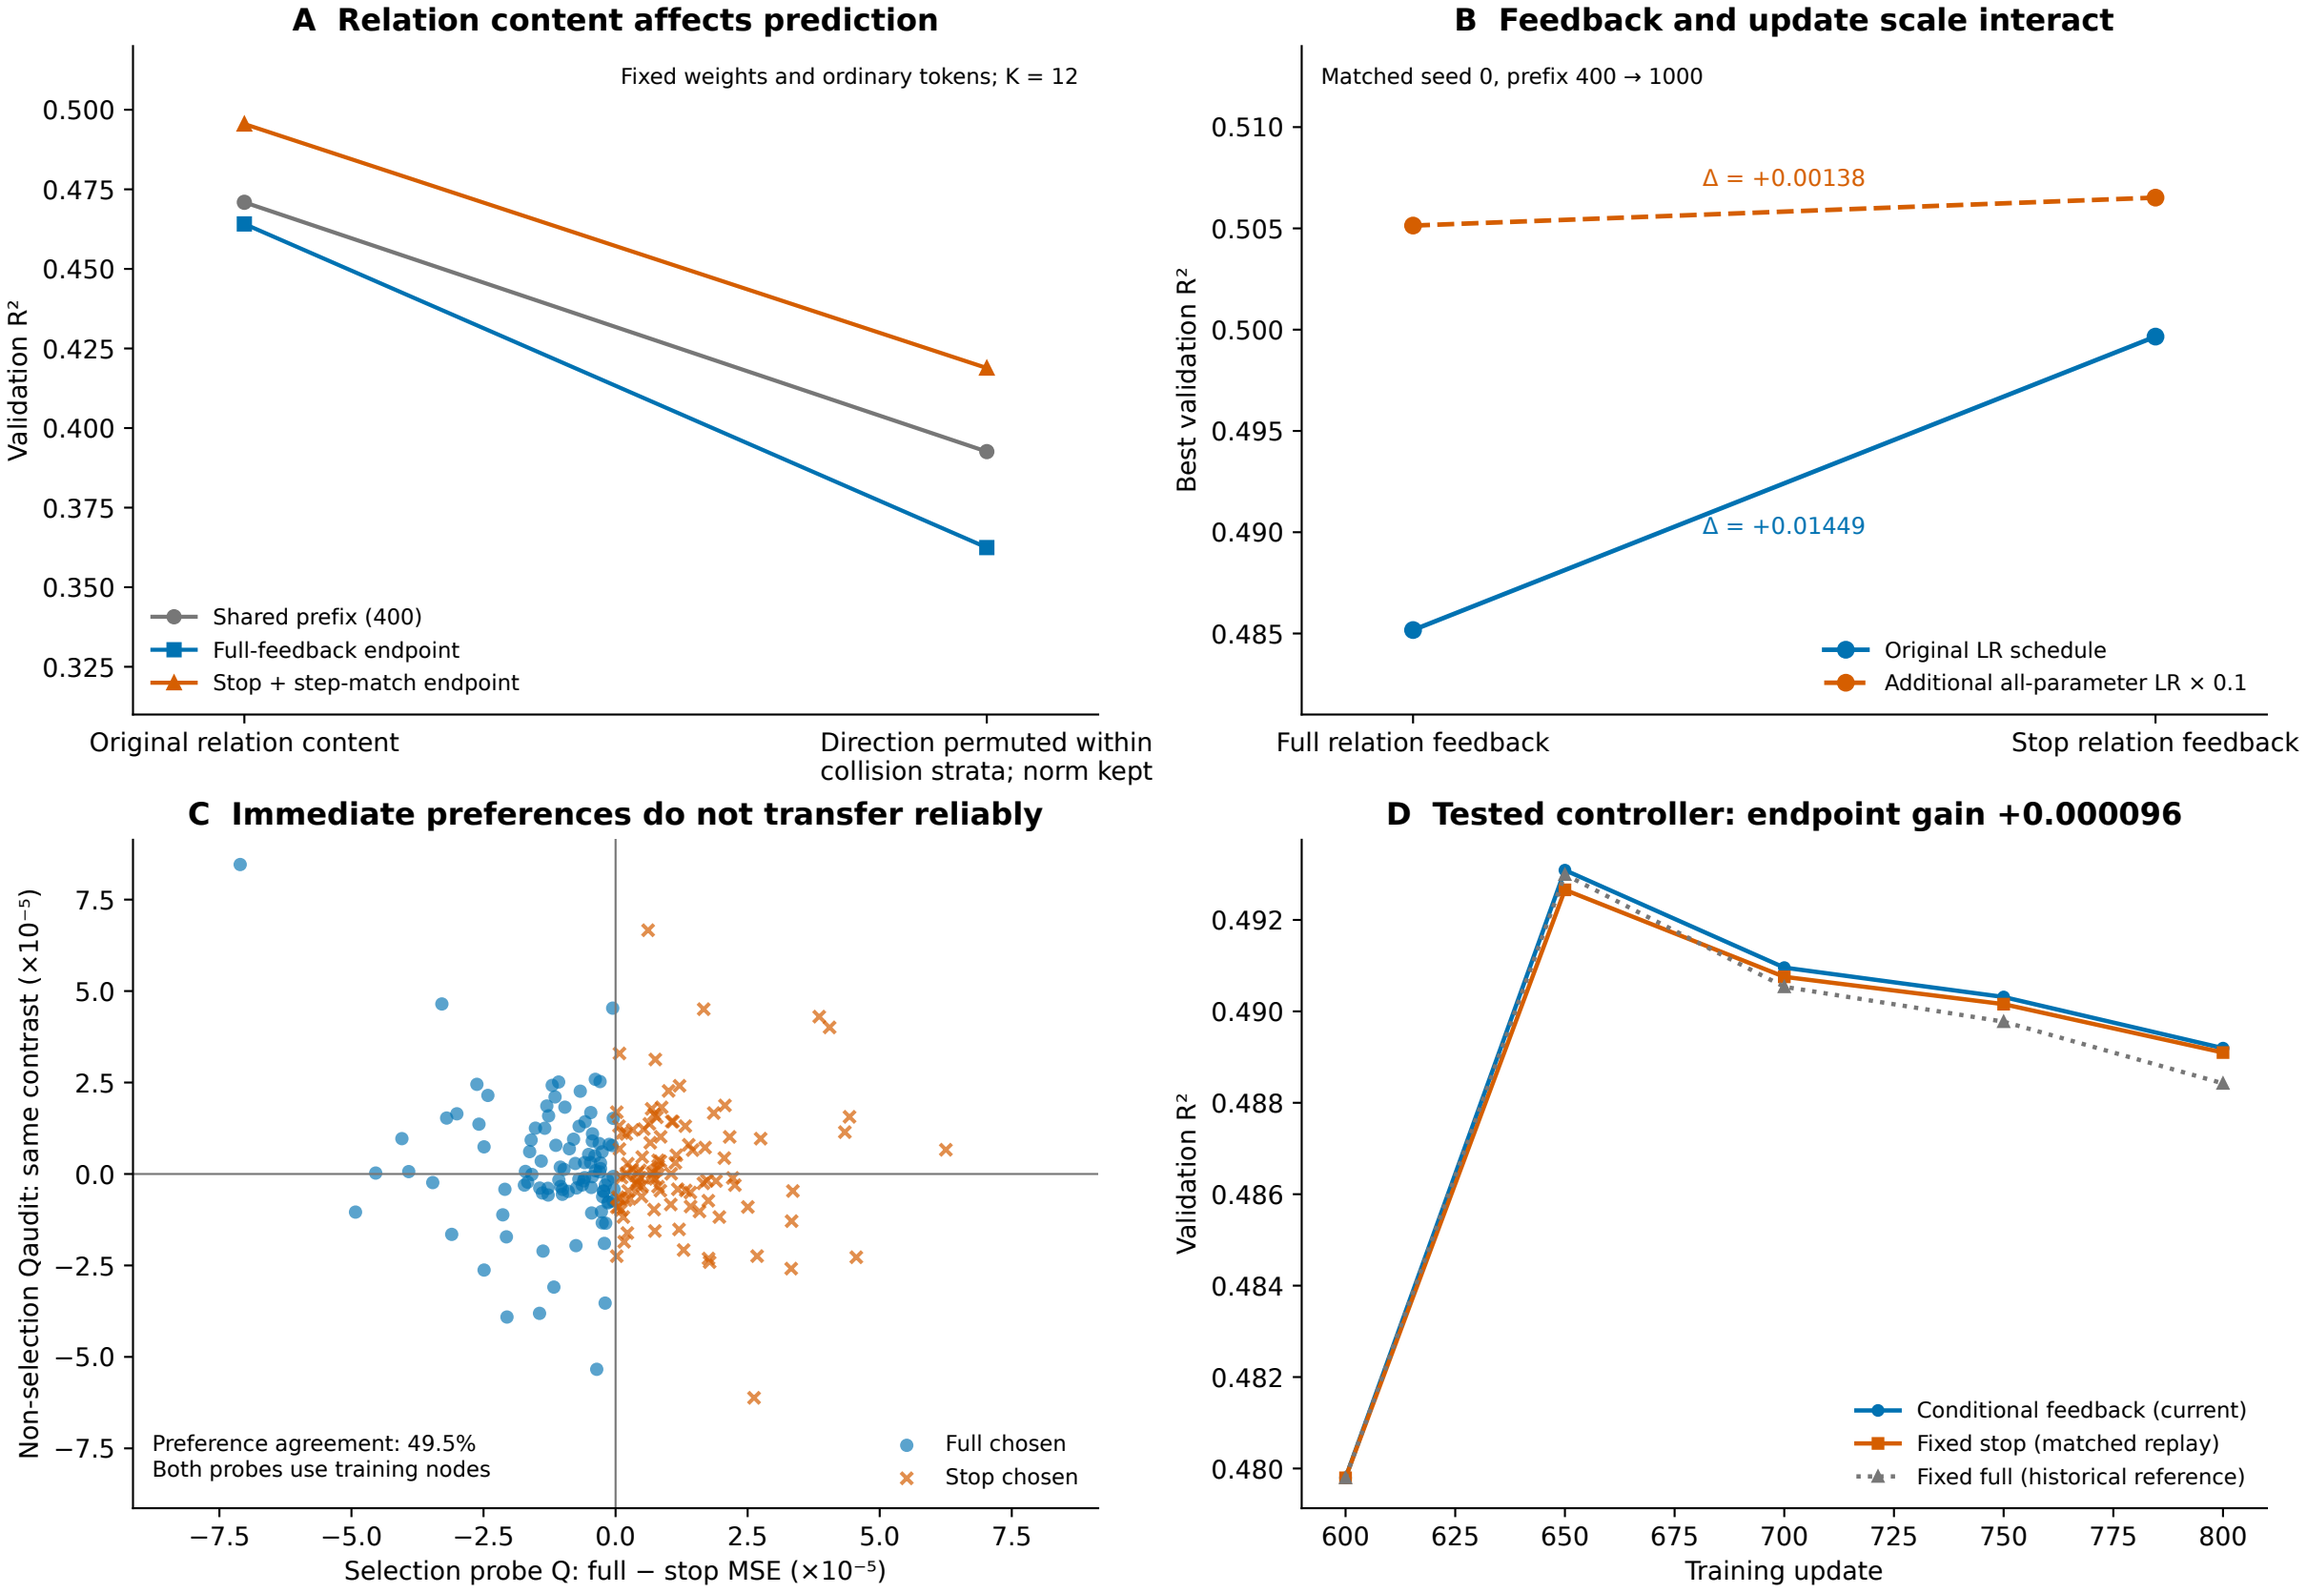
A Relation content affects prediction
Fixed weights and ordinary tokens; K = 12
Validation R²
0.500
0.475
0.450
0.425
0.400
0.375
0.350
0.325
Shared prefix (400)
Full-feedback endpoint
Stop + step-match endpoint
Original relation content
Direction permuted within collision strata; norm kept
B Feedback and update scale interact
Matched seed 0, prefix 400 → 1000
Best validation R²
0.510
0.505
0.500
0.495
0.490
0.485
Δ = +0.00138
Δ = +0.01449
Original LR schedule
Additional all-parameter LR × 0.1
Full relation feedback
Stop relation feedback
C Immediate preferences do not transfer reliably
Non-selection Qaudit: same contrast (×10⁻⁵)
7.5
5.0
2.5
0.0
−2.5
−5.0
−7.5
Preference agreement: 49.5%
Both probes use training nodes
Full chosen
Stop chosen
−7.5
−5.0
−2.5
0.0
2.5
5.0
7.5
Selection probe Q: full − stop MSE (×10⁻⁵)
D Tested controller: endpoint gain +0.000096
Validation R²
0.492
0.490
0.488
0.486
0.484
0.482
0.480
Conditional feedback (current)
Fixed stop (matched replay)
Fixed full (historical reference)
600
625
650
675
700
725
750
775
800
Training update